

ilifu Online Training

Robin Hall

User Training Workshop – Advanced Training

14 September 2021



Topics



- Python virtual environments
- R and RStudio
- Building Singularity containers

Python Virtual Environments

virtualenv

- Available anywhere
- <https://virtualenv.pypa.io/en/latest/>
- Isolated Python environment
- Less risk of conflicts occurring with `pip install --user`
- Similar to `venv` (`python -m venv`)
- Can customize which os python is used: `python2.7`, `python3+`
- Limited by os libraries

Python Virtual Environments

```
virtualenv --help
```

```
virtualenv /path/to/virtual_environment
```

```
--python
```

The Python interpreter to use

```
--system-site-packages
```

*Give the virtual environment
access to the global site-packages*

Python Virtual Environments

```
virtualenv /path/to/virtual_environment
```

Example:

```
virtualenv ~/.jupenv  
source ~/.jupenv/bin/activate  
(.jupenv) robh@slurm-login:~$  
which python  
    ~/.jupenv/bin/bin/python  
pip install scikit-learn  
deactivate (to exit)
```

Python Virtual Environments

Python virtualenv as a Jupyter kernel

Once the virtual environment is active:

```
ipython kernel install --name "<kernel_name>" --user
```

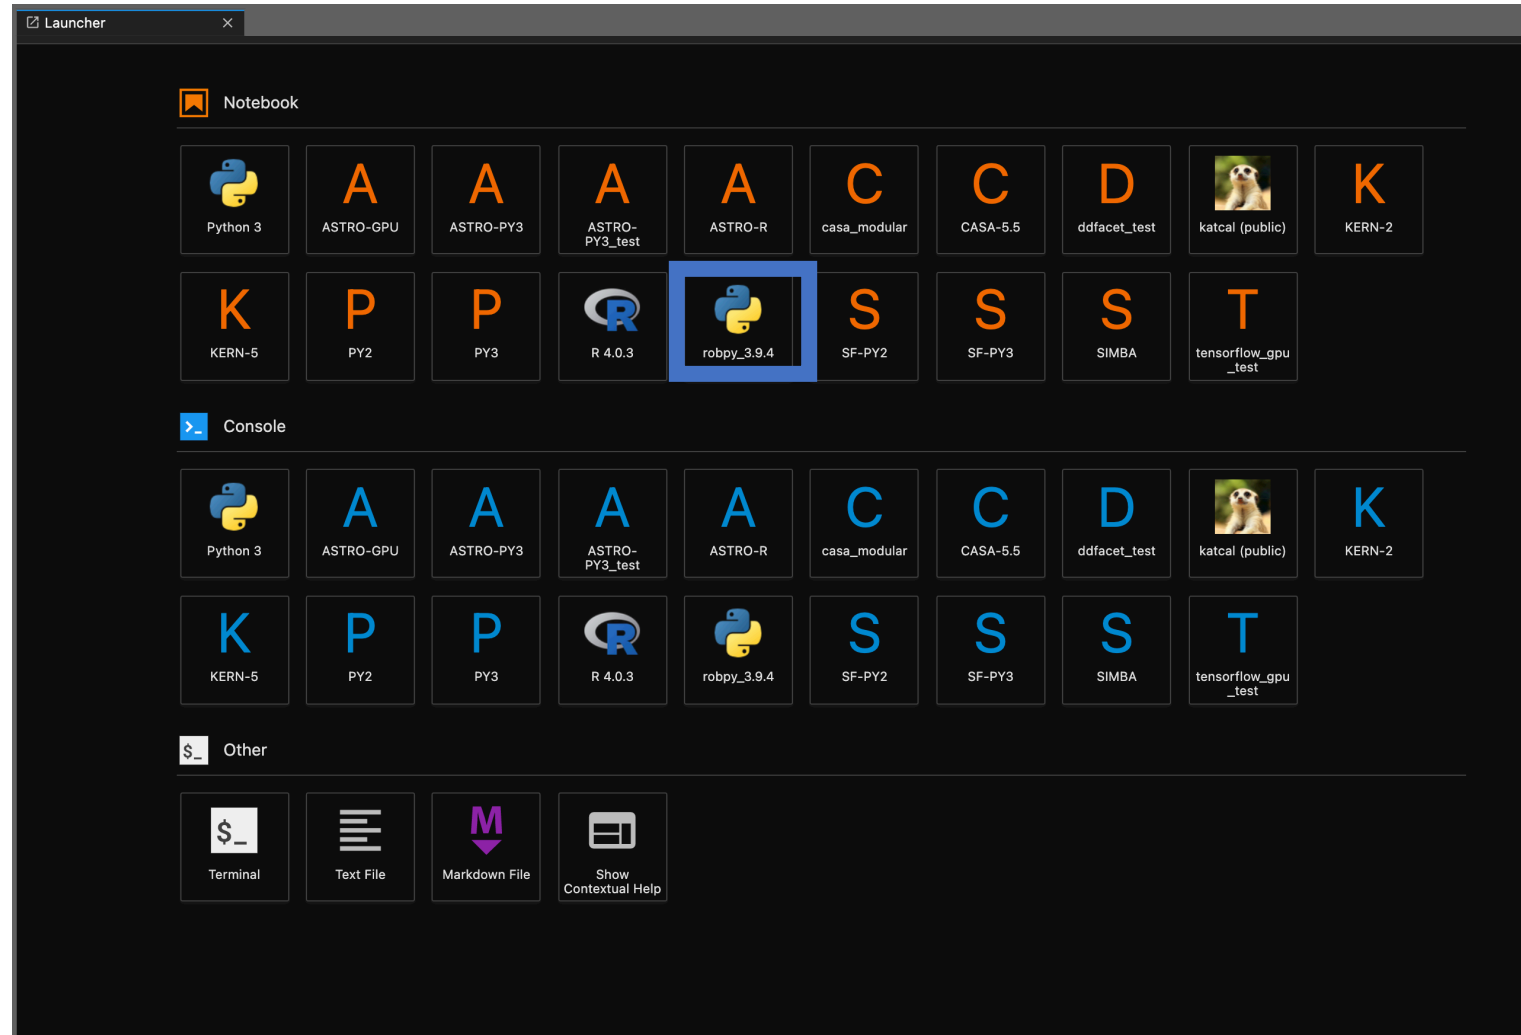
Example:

```
source ~/.jupenv/bin/activate  
pip install jupyter  
ipython kernel install --name "robpy_3.9.4" --user  
Installed kernelspec robpy_3.9.4 in /users/robh/.local/share/  
jupyter/kernels/robpy_3.9.4
```

Creates the kernel.json file at:

```
/users/robh/.local/share/jupyter/kernels/robpy_3.9.4/kernel.json
```

Python Virtual Environments



R and RStudio with slurm

http://docs.ilifu.ac.za/#/tech_docs/software_environments?id=running-r

When logged in via ssh:

```
robh@slurm-login:~$ srun --nodes=1 --tasks=1 --mem=8g --time 08:00:00  
--job-name="rstudio test" --pty bash
```

```
robh@compute-004:~$ module add R/RStudio1.2.5042-R4.0.0
```

```
robh@compute-004:~$ rstudio
```

The environment variable RSTUDIO_PASSWORD was not set, so your password has been chosen for you. It's: *****

Running rserver on port 40755

To connect to this server run this on your local machine:

```
ssh -A robh@compute-004 -o "ProxyCommand=ssh  
robh@slurm.ilifu.ac.za nc compute-004 22"  
-L8081:localhost:40755
```

then visit <http://localhost:8081> in your browser and use the username "robh" to login with the password *****

R and Studio with slurm

http://docs.ilifu.ac.za/#/tech_docs/software_environments?id=running-r

On your local machine:

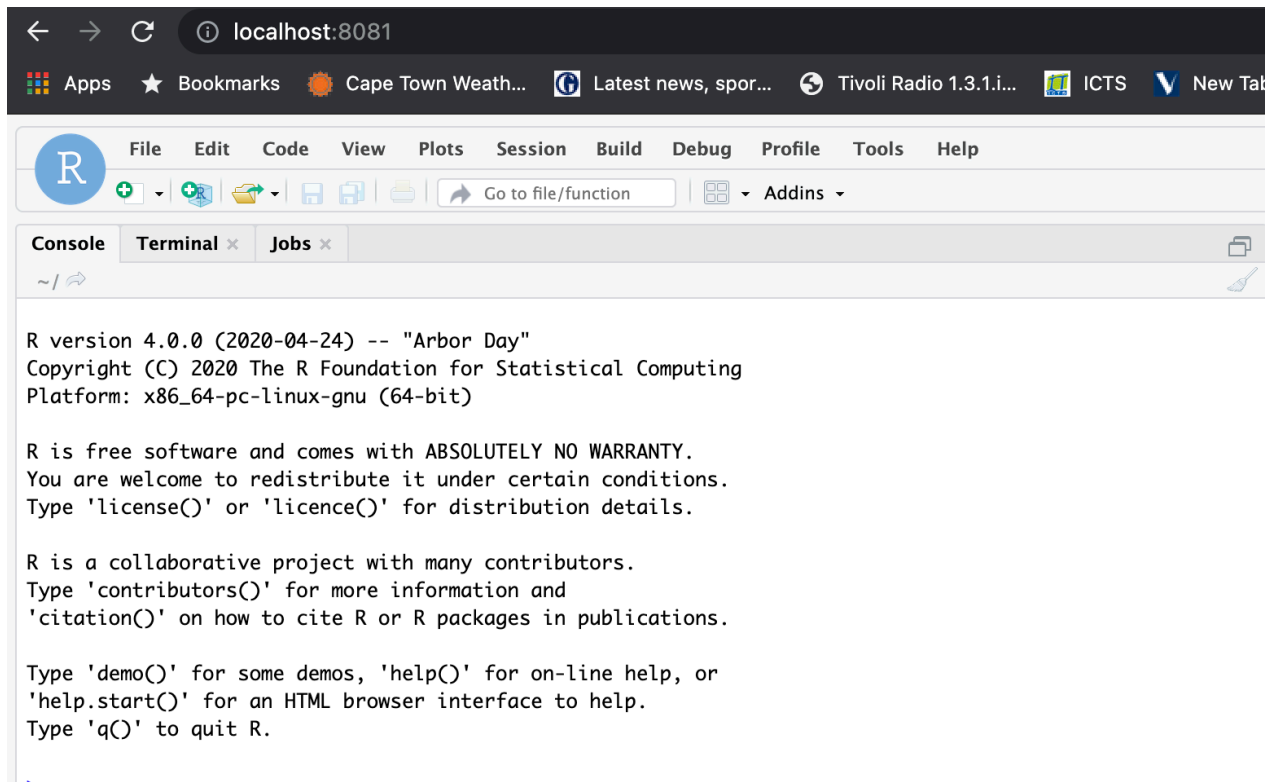
```
robinh@MacBook-Pro[~] -> % ssh -A robh@compute-004 -o  
"ProxyCommand=ssh robh@slurm.ilifu.ac.za nc  
compute-004 22" -L8081:localhost:40755
```

Go to: <http://localhost:8081>

R and Studio with slurm

http://docs.ilifu.ac.za/#/tech_docs/software_environments?id=running-r

Now you can access RStudio through your web browser:



Singularity containers

Singularity 3.0 definition files

https://sylabs.io/guides/3.0/user-guide/definition_files.html



Singularity containers

```
Bootstrap: docker
From: ubuntu:18.04
```

Header:

- base OS
- Keywords depend on selection

%files

```
/file1
/file1 /opt
```

%environment

```
export LISTEN_PORT=12345
export PYTHONPATH=training_demo/:$PYTHONPATH
```

%post

```
apt update
apt-get install -y python3 python3-pip git
```

%runscript

```
python3 "$@"
```

Singularity containers

Bootstrap: docker
From: ubuntu:18.04

%files

```
/file1  
/file1 /opt
```



Files:

- Copy files from host to container
- Source and destination

%environment

```
export LISTEN_PORT=12345  
export PYTHONPATH=training_demo/:$PYTHONPATH
```

%post

```
apt update  
apt-get install -y python3 python3-pip git
```

%runscript

```
python3 "$@"
```

Singularity containers

```
Bootstrap: docker
From: ubuntu:18.04
```

%files

```
/file1
/file1 /opt
```

%environment

```
export LISTEN_PORT=12345
export PYTHONPATH=training_demo/:$PYTHONPATH
```

Environment:

- Variables set at runtime

%post

```
apt update
apt-get install -y python3 python3-pip git
```

%runscript

```
python3 "$@"
```

Singularity containers

Bootstrap: docker
From: ubuntu:18.04

%files

```
/file1  
/file1 /opt
```

%environment

```
export LISTEN_PORT=12345  
export PYTHONPATH=training_demo/:$PYTHONPATH
```

%post

```
apt update  
apt-get install -y python3 python3-pip git
```

%runscript

```
python3 "$@"
```

Post:

- Commands to build on OS
- Customise container

Singularity containers

Bootstrap: docker

From: ubuntu:18.04

%files

```
/file1  
/file1 /opt
```

%environment

```
export LISTEN_PORT=12345  
export PYTHONPATH=training_demo/:$PYTHONPATH
```

%post

```
apt update  
apt-get install -y python3 python3-pip git
```

%runscript

```
python3 "$@"
```

Runscript:

- Executed when container is run

Singularity containers

Build container:

```
sudo singularity build <definition_file.def>  
<container_name.simg>
```

Run container:

```
singularity run <container_name.simg> <args>
```

```
singularity exec <container_name.simg> <args>
```

```
singularity shell <container_name.simg>
```



Demo time



Summary of Software Environs.

virtual environment

- Good for prototyping and rapid development
- Can be used by a group but needs to be in appropriate folder

Module

- Useful if container doesn't have version of software
- Some modules execute containers more conveniently

Containers

- Best for reproducibility and sharing
- Can be used by anyone with the path