

# ilifu Online Training

Robin Hall

User Training Workshop – Basic Training

7 September 2021

# Topics

- Introduction to the ilifu research facility services
- Directory structure
- Software environment
- Accessing the ilifu services
- Using JupyterHub

# High Performance Computing

Combining power of distributed computers

- Collection of servers (computers)
- Connected by fast local network

Some terminology

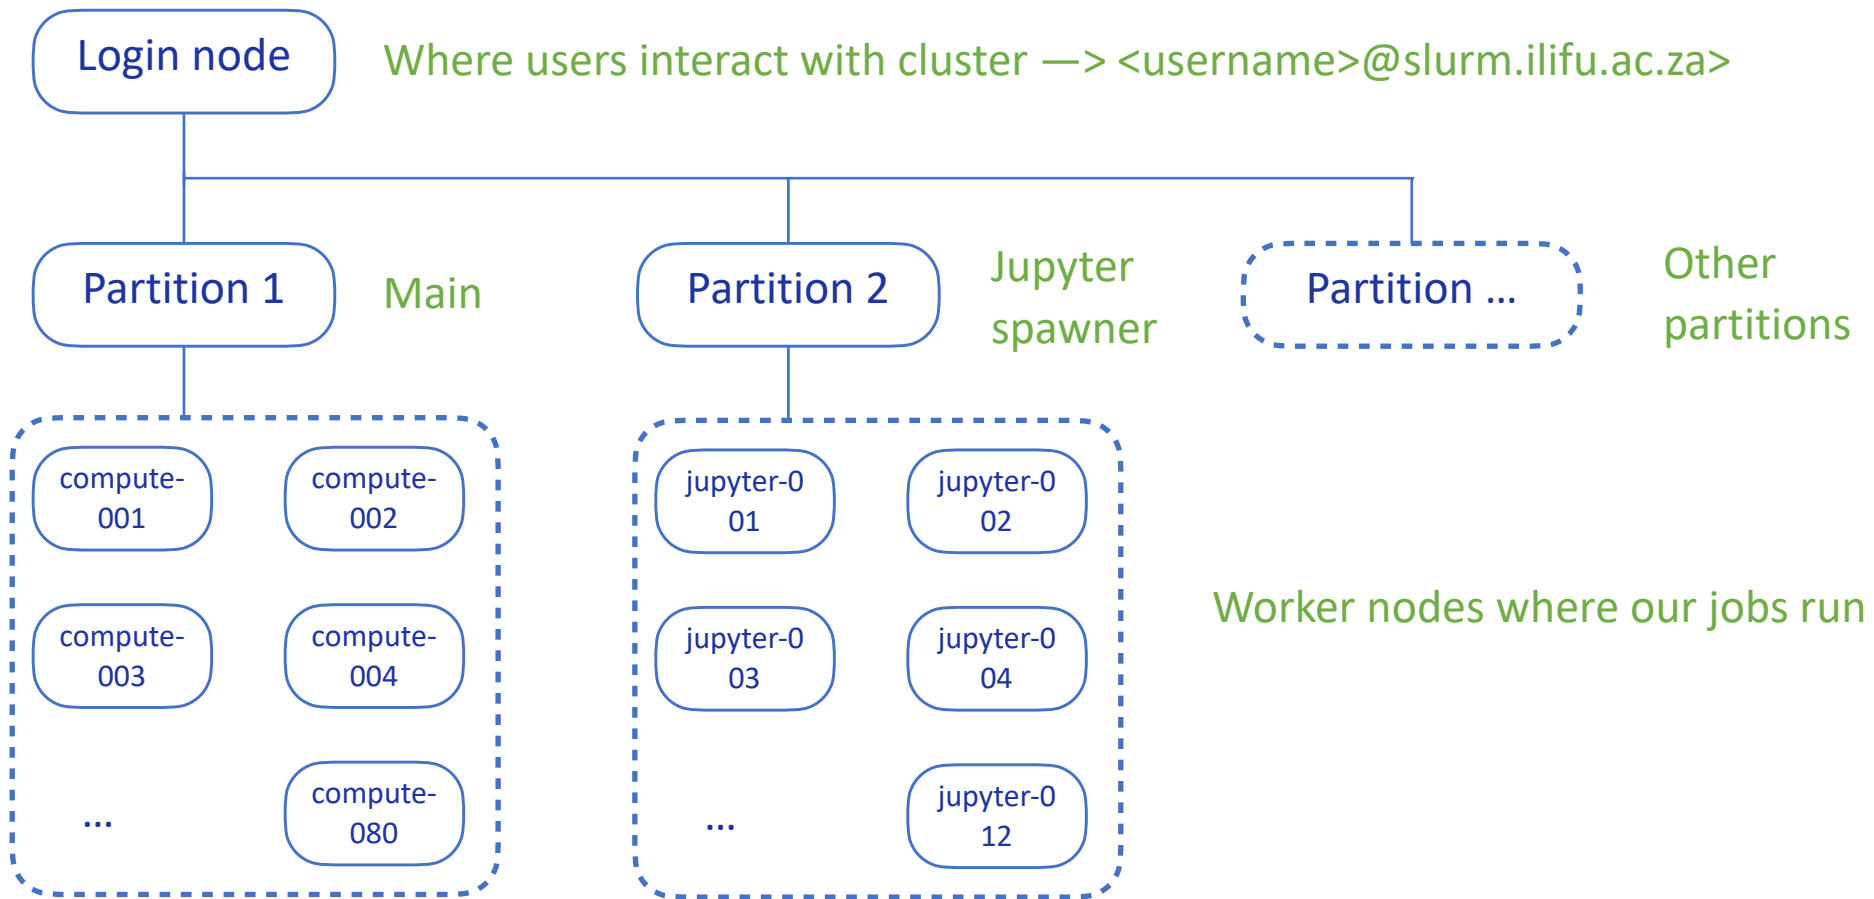
- Servers also referred to as nodes
- Group of nodes is a cluster

# ilifu Research Facility

## Cloud-based infrastructure for data-intensive research

- Support variety of different scientific projects and requirements
- Flexible compute environment
- Data management: storage, transfer

# ilifu Research Facility



## Software we use

- Job Scheduler to manage resources - Slurm



- JupyterHub service – development environ.



- Containerised software environment - Singularity



- Other services: data transfer, CARTA

# Computing environment - interface

## ssh – shell terminal

```
* Support: https://ubuntu.com/advantage

System information as of Fri Aug 23 11:36:57 SAST 2019

System load: 0.49      Users logged in: 8
Usage of /: 35.9% of 21.15GB IP address for ens3: 192.168.100.39
Memory usage: 5%      IP address for ens4: 10.102.26.97
Swap usage: 0%        IP address for ens5: 10.102.28.133
Processes: 396

* Keen to learn Istio? It's included in the single-package MicroK8s.

https://snapcraft.io/microk8s

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

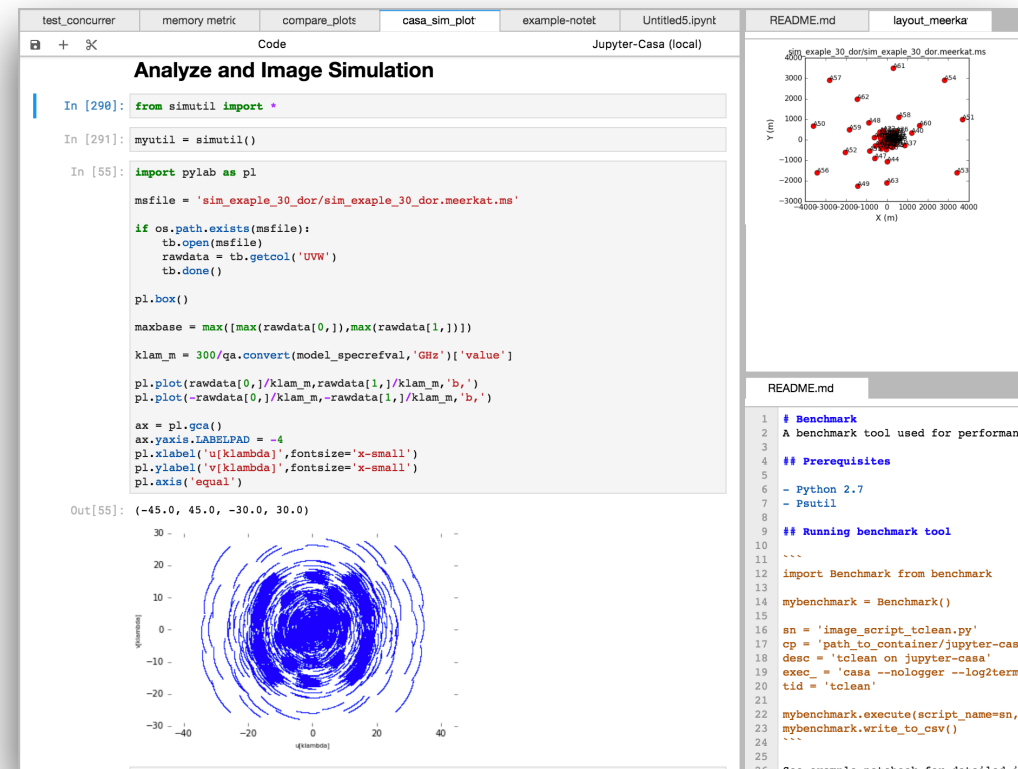
* Canonical Livepatch is available for installation.
- Reduce system reboots and improve kernel security. Activate at:
https://ubuntu.com/livepatch

170 packages can be updated.
75 updates are security updates.

Last login: Fri Aug 23 09:08:21 2019 from 196.11.235.232
[jeremy@slurm-login:~$ sinfo
PARTITION AVAIL TIMELIMIT NODES STATE NODELIST
Main*      up 14-00:00:0 8 mix slwrk-[106-113]
Main*      up 14-00:00:0 14 alloc slwrk-[101,104-105,114-124]
Main*      up 14-00:00:0 38 idle slwrk-[102-103,125-160]
JupyterSpawnerONLY up infinite 4 mix slwrk-[201-202,205,209]
JupyterSpawnerONLY up infinite 4 alloc slwrk-[206-208,210]
JupyterSpawnerONLY up infinite 2 idle slwrk-[203-204]
jeremy@slurm-login:~$ sbatch compute_job.sh]
```

ssh <username>@slurm.ilifu.ac.za

## JupyterHub



<https://jupyter.ilifu.ac.za>

# Computing environment - ssh

## Your SSH key

- Used in the SSH (Secure Shell) protocol
- Authentication method for gaining access to encrypted connection between systems
- Use connection to manage system remotely
- We need your SSH key so our system knows to let you in



# Computing environment - ssh

## Generating SSH key

- If you don't already have one
- New computer

GitHub Docs on key generation:

<https://docs.github.com/en/github/authenticating-to-github/connecting-to-github-with-ssh/generating-a-new-ssh-key-and-adding-it-to-the-ssh-agent>

# Directory Structure

- Common areas:
  - /users
    - only 40 TB shared among all users, for scripts and small files – **don't place data here, capping /users storage capacity can prevent access to the cluster for all users.**
  - /scratch3/users
    - storage space for processing data, temporary storage only, i.e. use this space during processing, and then clear all files immediately after processing. Remove unnecessary data and move data that you want to keep to project folder. /scratch3 has recently been increased from 100 TB to 400 TB
- Remaining storage is separated by group:
  - IDIA, CBIO, ILIFU

# Directory Structure

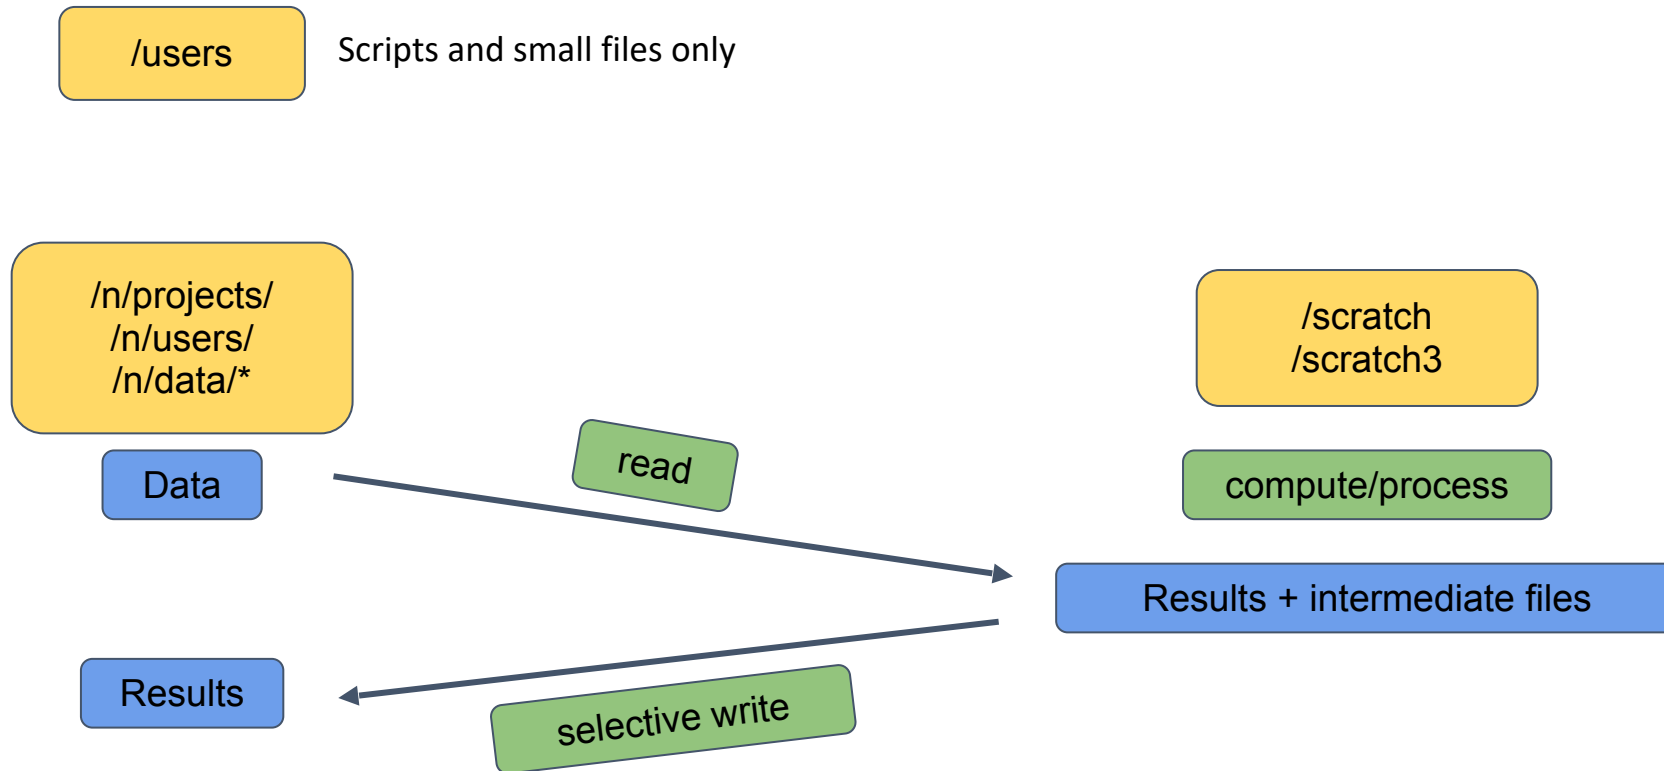
## IDIA structure:

- /idia/users
  - user's private work directory, may store data products that are not ready to move to shared project space
- /idia/projects
  - project specific directories. These directories are for sharing data and resources within project groups. Raw data associated with a project will also be available from the project folder. Raw data should always be read-only.
- /idia/software
  - software containers and the IDIA Pipelines software is stored here
- /idia/software/containers

# Directory Structure

- Similar structure for /cbio and /ilifu groups,
  - /cbio/users
  - /cbio/projects
  - /cbio/soft
- /ilifu/users
- /ilifu/software
- Exception for ilifu projects:
  - /ilifu/astro/projects
  - /ilifu/bio/projects

# Directory Structure - workflow



\*/n/data generally read-only

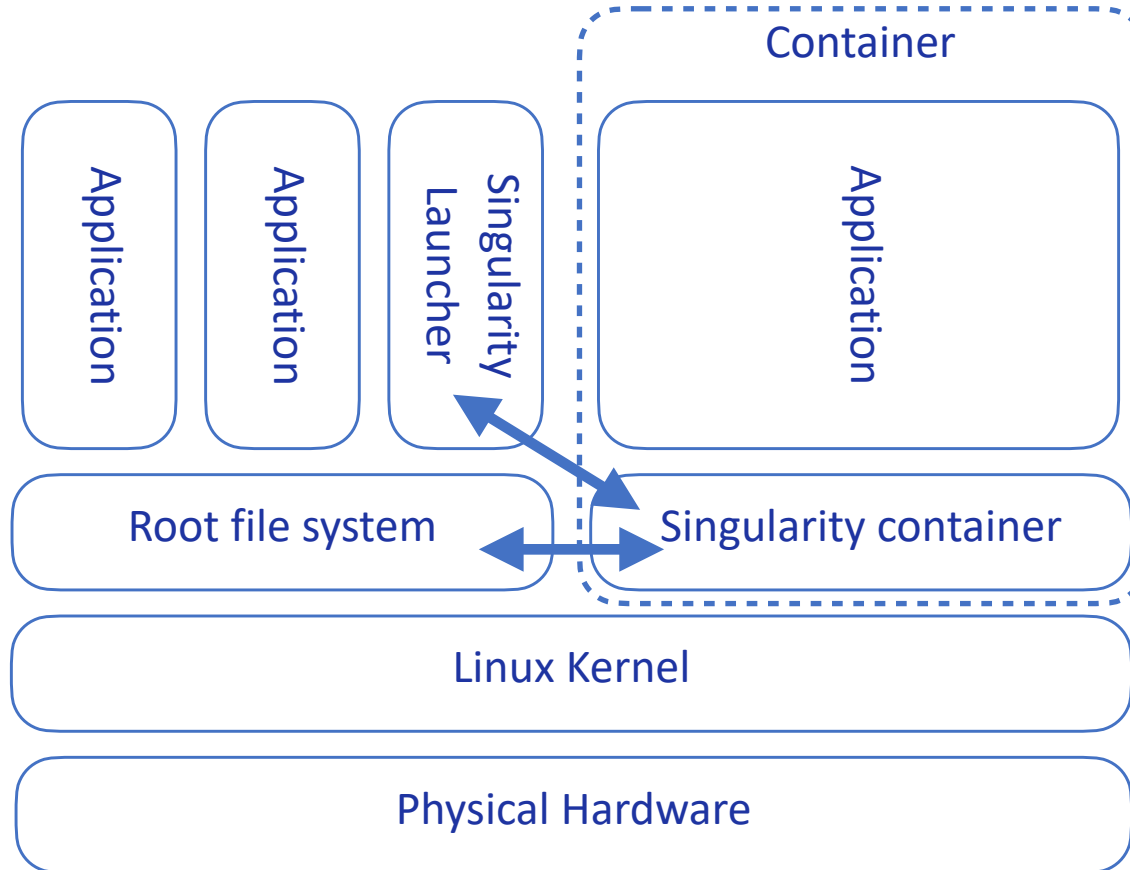
# Software environment - Singularity containers



- Encapsulated software environments
- A software stack that contains everything required to run an application/workflow, including files, environmental variables, libraries and dependencies
- Containers accessible across platforms and services, allowing sharing of application environments



# Software environment - Singularity containers



# Software environment - Singularity containers



## Supported Containers:

- CASA
  - Astronomy container (ASTRO-PY3)
  - KERN suite
  - Python 2.7, Python 3.6, R
  - GPU container
  - Project containers:
    - MeerLICHT,
    - LADUMA
    - HI intensity map
- 
- `/idia/software/containers`
  - `/ilifu/software/containers`



# Software environment - Singularity containers



**Open container as an interactive shell:**

```
singularity shell /path/to/container
```

**Example:**

```
$ singularity shell /idia/software/containers/ASTRO-PY3.simg
```

**Run a script/workflow using a container environment:**

```
singularity exec /path/to/container <software> <script/input_params>
```

```
$ singularity exec /idia/software/containers/casa-stable.img casa -c myscript.py
```

# Software environment - modules

## module avail

```
~$ module avail
```

```
----- /software/modules/common -----
LAPACK/3.9.0
2.6.6
R/RStudio1.2.5042-R4.0.0
R/RStudio1.2.5042-R4.0.4
R/3.6.3
R/4.0.0
R/4.0.2
R/4.0.3 (D)
anaconda3/login
anaconda3/login
anaconda3/2020.07 (D)
cuda/10.0.130_410.48
cuda/10.1.243_418.87.00
cuda/10.2.89_440.33.01
cuda/11.0.2_450.51.05 (D)
drmaa/1.1.1
homebrew/2.4.13
java/jre-1.8.0_261
java/openjdk-14.0.1 (D)
julia/1.5.3
maven/3.6.3
mono/6.8.0.123
openBLAS/0.3.9
openmpi/2.1.6
openmpi/3.1.6
openmpi/4.0.3
openmpi/4.0.5
openmpi/4.1.0 (D)
perl/5.33.0
perlbrew/perlbrew
python/2.7.18
python/3.7.7
python/3.8.2
python/3.8.3
python/3.8.6
python/3.9.0 (D)
ruby/

----- /software/modules/astro -----
casa/5.7.0
casa/5.7.2-4
casa/6.1.0-118-monolithic
casa/6.1.2.7-pipeline
casa/6.1.2.7-modular (D)
pybdsf/1.9.2

----- /software/modules/bio -----
bcbio/1.2.3
bcftools/1.10.2
canvas/1.40.0.1613
gemini/gemini
genomestrip/2.00.1958
htslib/1.10.2
plink/2.00a2.3
popgen/0.1
prsize-2/2.3.1d
samtools/1.10
treePL/homebrew
vcftools/0.1.16
vep/singularity
vep/101.0 (D)

----- /usr/share/lmod/lmod/modulefiles -----
Core/lmod/6.6
Core/settarg/6.6
```

Where:

D: Default Module

# Software environment - modules

- module avail
- module help <module>

```
~$ module help python
```

```
----- Module Specific Help for "python/3.9.0" -----  
This module configures Python 3.9.0 for use
```

- module load <module>
- module purge
- module --help

# ilifu

## JupyterHub

- <https://jupyter.ilifu.ac.za>



# ilifu jupyter login

Sign in

**Username:**

**Password:**

Sign In

- Session size

## Server Options

### Nodes Free

as at Wed Aug 25 14:18:01 SAST 2021

21 Minimum

9 Small

3 Medium

1 Large

0 Half-Max

0 Max

4 GPU

Select a job profile:

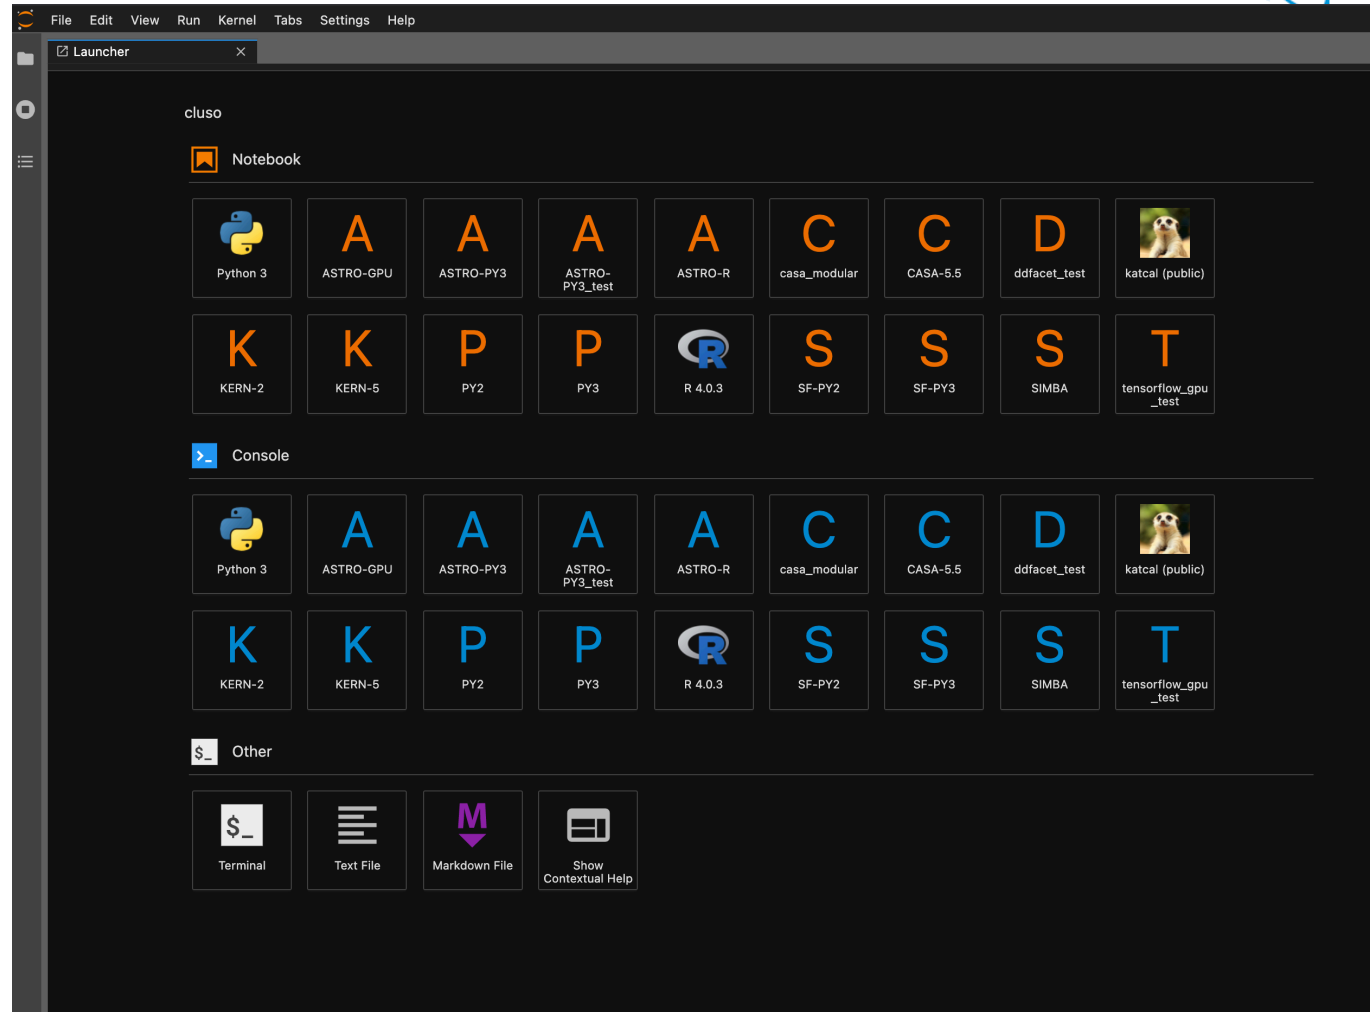
Minimum Node - 1 core, 7 GB, 18 hours idle timeout, max 5 days lifespan



Start

# JupyterHub

- Choose kernel in launcher



# Getting help

- Support contact  
[support@ilifu.ac.za](mailto:support@ilifu.ac.za)
- User documentation  
<http://docs.ilifu.ac.za/#/>
- Ilifu System Status  
<https://status.ilifu.ac.za/>
- Training videos  
<http://www.ilifu.ac.za/il/accessing-facilities/training>



Demo time

